

Breeze LES master report

Data/status checked: 2026-09-20T01:47:54.123Z

DYCOMS: 15 completed cases. GABLS1: 12 of 15 planned cases have completed, audited local results.

This document combines the completed DYCOMS RF01 experiment (including the Pressel comparison) with the developing GABLS1 experiment. The GABLS run order is smallest grids first; incomplete runs contribute status information only, never scientific curves.

Read finite-output completion separately from physical fidelity: all three unbounded DYCOMS Centered2 cases developed negative total water. The DYCOMS chapter retains that audit.

GABLS1 uses a fixed 400 m cube and 9 h of surface cooling and geostrophic forcing. Its no-interior-closure configurations retain MOST wall stress and heat exchange. The question is whether numerical choices preserve resolved shear-driven turbulence on coarser grids.

Contents: Part I - completed DYCOMS and Pressel comparison; Part II - GABLS setup, progress, numerical references, then audited results as they arrive.

Part III adds GABLS3: an observational evaluation now in preparation. All prior DYCOMS and GABLS1 information is retained. The requested persistent evaluation repository is [glwagner/BreezeEvaluation.jl](https://github.com/glwagner/BreezeEvaluation.jl).

Breeze DYCOMS-II RF01: resolution and numerics

Complete: 15 cases analyzed

Data checked: 2026-09-19T11:58:17.096Z

All 15 requested four-hour cases are analyzed. This matrix does not establish faithful low-resolution DYCOMS turbulence: the four coarse WENO cases retain the wrong third-moment sign in the 600–760 m layer, and refining WENO9/no-SGS from 35 m to 10 m horizontal spacing substantially improves the third moment despite little improvement in variance-marker error. Fine WENO9/no-SGS has the smallest third-moment marker RMSE among the tested WENO cases. All three unbounded Centered2 cases develop negative total water and cannot establish physical fidelity. Cloud persistence alone is therefore insufficient. These are single-seed comparisons; the fine grid is not a proven converged solution, and the larger coarse domain confounds the resolution comparison.

Completed Centered2 runs contain physically inadmissible negative total water (sampled minima: coarse: -4.338 g/kg; canonical: -7.697 g/kg; fine: -7.401 g/kg). Their cloud and turbulence statistics are retained as diagnostics of this unbounded numerical configuration, not evidence of a physically faithful solution. The production extrema figure shows the full sampled time history; verification here means complete finite output, not physical validity.

Centered2 smoke audit: all three grids stayed finite over the first 600 simulated seconds, but canonical and fine cases produced negative total water. Across the 11 one-minute samples, minimum q_t was coarse: 1.0872 g/kg; canonical: -0.0605 g/kg; fine: -1.3679 g/kg. Passing the finite-value gate establishes neither moisture positivity nor physical validity. Production retains unbounded centered moisture; full-run extrema are reported separately in the production audit. These are smoke-test results, not completed four-hour Centered2 statistics.

At 10 m horizontal / 5 m vertical spacing with WENO9, adding Smagorinsky changes fourth-hour minimum w'^3 over 600–760 m from -0.0622 to -0.0304 m^3/s^3 , third-moment marker RMSE from 0.0159 to 0.0313 m^3/s^3 , and variance-marker RMSE from 0.0631 to 0.0769 m^2/s^2 . LWP changes from 55.48 to 53.23 g/m^2 . Solver wall increases by 10.1%. This closure weakens the negative third moment and increases both marker errors in this paired single-realization comparison; it is not evidence against every possible closure or coefficient choice.

Breeze DYCOMS-II RF01: resolution and numerics

At 10 m horizontal / 5 m vertical resolution without SGS, WENO5 uses 70.5 solver wall minutes versus 92.8 for WENO9 (24.0% less). However, its fourth-hour third-moment RMSE against digitized in-situ markers is 0.0310 rather than 0.0159 m^3/s^3 . Its minimum w'^3 in 600–760 m is -0.0372 versus -0.0622 m^3/s^3 . Lower advection order saves wall time but weakens the negative third moment in this realization. Timings include output within the solve and exclude startup/compilation before the solve; they are not isolated kernel benchmarks.

Fine Centered2 with Smagorinsky completes in 70.2 solver wall minutes, 31.3% less than fine WENO9 with Smagorinsky. Its fourth-hour LWP is 44.17 versus 53.23 g/m^2 , and third-moment marker RMSE is 0.0382 versus 0.0313 m^3/s^3 . Negative total water prevents interpreting this cheaper run as an acceptable physical solution. Timing includes diagnostics and differing timestep counts; no GPU profiling was performed to identify the cost mechanism.

At fixed 5 m vertical spacing and fixed horizontal domain, changing WENO9/no-SGS from 35 m to 10 m horizontal spacing changes the fourth-hour minimum w'^3 over 600–760 m from -0.0126 to -0.0622 m^3/s^3 . The fine case is negative throughout 600–760 m, agreeing with the sign of the retained in-situ observations there.

For this same 35 m to 10 m comparison, unweighted RMSE against the digitized in-situ third-moment markers changes from 0.0454 to 0.0159 m^3/s^3 ; variance-marker RMSE changes from 0.0606 to 0.0631 m^2/s^2 . The third moment improves substantially, while variance does not improve by this metric. These are single realizations without observational error-bar weighting, not a convergence demonstration.

All 4 completed coarse WENO cases have positive fourth-hour w'^3 throughout 600–760 m, while the retained in-situ observations there are negative. Cloud persistence therefore does not establish faithful coarse-grid turbulence.

Coarse, WENO9 without closure: fourth-hour LWP 49.87 g/m^2 , cloud fraction 99.34%, peak w^2 0.395 m^2/s^2 , and Δq : 0.019 g/kg .

At coarse resolution without closure, WENO9 minus WENO5 changes fourth-hour LWP by 3.42 g/m^2 and peak w^2 by 0.099 m^2/s^2 . These are paired single-realization differences, not statistically established scheme effects.

At coarse resolution with WENO9, adding Smagorinsky changes fourth-hour LWP by -6.54 g/m^2 and peak w^2 by -0.038 m^2/s^2 . Third-moment marker RMSE changes from 0.0704 to 0.0595 m^3/s^3 .

Breeze DYCOMS-II RF01: resolution and numerics

Canonical, WENO9 without closure: fourth-hour LWP 61.69 g/m², cloud fraction 99.91%, peak w^2 0.521 m²/s², and Δq_t 0.003 g/kg.

At canonical resolution without closure, WENO9 minus WENO5 changes fourth-hour LWP by 2.09 g/m² and peak w^2 by 0.054 m²/s². These are paired single-realization differences, not statistically established scheme effects.

At canonical resolution with WENO9, adding Smagorinsky changes fourth-hour LWP by -9.50 g/m² and peak w^2 by -0.067 m²/s². Third-moment marker RMSE changes from 0.0454 to 0.0444 m³/s³.

Fine, WENO9 without closure: fourth-hour LWP 55.48 g/m², cloud fraction 99.95%, peak w^2 0.472 m²/s², and Δq_t 0.031 g/kg.

At fine resolution without closure, WENO9 minus WENO5 changes fourth-hour LWP by -2.06 g/m² and peak w^2 by -0.011 m²/s². These are paired single-realization differences, not statistically established scheme effects.

At fine resolution with WENO9, adding Smagorinsky changes fourth-hour LWP by -2.25 g/m² and peak w^2 by -0.027 m²/s². Third-moment marker RMSE changes from 0.0159 to 0.0313 m³/s³.

Coarse, Centered2 with Smagorinsky: sampled global q_t extrema over 0–4 h were -4.3385 to 15.0839 g/kg; 227 of 241 samples contained negative moisture. These are 60-second samples, not bounds on every timestep.

At coarse resolution with Smagorinsky, Centered2 minus WENO9 changes fourth-hour LWP by 5.52 g/m² and peak w^2 by 0.053 m²/s². Centered moisture is unbounded; this changes scheme family and bounding as well as formal order.

At coarse resolution with Smagorinsky, Centered2 minus WENO5 changes fourth-hour LWP by 6.59 g/m² and peak w^2 by 0.095 m²/s². Centered moisture is unbounded; this changes scheme family and bounding as well as formal order.

Canonical, Centered2 with Smagorinsky: sampled global q_t extrema over 0–4 h were -7.6972 to 17.4041 g/kg; 231 of 241 samples contained negative moisture. These are 60-second samples, not bounds on every timestep.

At canonical resolution with Smagorinsky, Centered2 minus WENO9 changes fourth-hour LWP by -16.49 g/m² and peak w^2 by -0.027 m²/s². Centered moisture is unbounded; this changes scheme family and bounding as well as formal order.

Breeze DYCOMS-II RF01: resolution and numerics

At canonical resolution with Smagorinsky, Centered2 minus WENO5 changes fourth-hour LWP by -15.33 g/m^2 and peak w^2 by $-0.006 \text{ m}^2/\text{s}^2$. Centered moisture is unbounded; this changes scheme family and bounding as well as formal order.

Fine, Centered2 with Smagorinsky: sampled global q_t extrema over 0–4 h were -7.4007 to 18.0746 g/kg ; 234 of 241 samples contained negative moisture. These are 60-second samples, not bounds on every timestep.

At fine resolution with Smagorinsky, Centered2 minus WENO9 changes fourth-hour LWP by -9.06 g/m^2 and peak w^2 by $-0.072 \text{ m}^2/\text{s}^2$. Centered moisture is unbounded; this changes scheme family and bounding as well as formal order.

At fine resolution with Smagorinsky, Centered2 minus WENO5 changes fourth-hour LWP by -7.54 g/m^2 and peak w^2 by $-0.082 \text{ m}^2/\text{s}^2$. Centered moisture is unbounded; this changes scheme family and bounding as well as formal order.

Fourth-hour metrics

Coarse / WENO9 / no SGS: LWP 49.87 g/m²; cloud 99.34%; peak w^2 0.395 m²/s²; Δq_t 0.019 g/kg; E 4.75 mm/s.

Coarse / WENO5 / no SGS: LWP 46.45 g/m²; cloud 98.91%; peak w^2 0.296 m²/s²; Δq_t 0.015 g/kg; E 4.70 mm/s.

Coarse / WENO9 / Smagorinsky: LWP 43.33 g/m²; cloud 99.07%; peak w^2 0.356 m²/s²; Δq_t 0.080 g/kg; E 5.21 mm/s.

Coarse / WENO5 / Smagorinsky: LWP 42.27 g/m²; cloud 98.52%; peak w^2 0.314 m²/s²; Δq_t 0.047 g/kg; E 4.93 mm/s.

Coarse / Centered2 / Smagorinsky: LWP 48.85 g/m²; cloud 99.44%; peak w^2 0.409 m²/s²; Δq_t 0.214 g/kg; E 7.27 mm/s.

Canonical / WENO9 / no SGS: LWP 61.69 g/m²; cloud 99.91%; peak w^2 0.521 m²/s²; Δq_t 0.003 g/kg; E 4.32 mm/s.

Canonical / WENO5 / no SGS: LWP 59.61 g/m²; cloud 99.92%; peak w^2 0.467 m²/s²; Δq_t 0.020 g/kg; E 5.19 mm/s.

Canonical / WENO9 / Smagorinsky: LWP 52.19 g/m²; cloud 99.37%; peak w^2 0.454 m²/s²; Δq_t 0.023 g/kg; E 5.24 mm/s.

Canonical / WENO5 / Smagorinsky: LWP 51.04 g/m²; cloud 98.95%; peak w^2 0.433 m²/s²; Δq_t 0.060 g/kg; E 5.24 mm/s.

Canonical / Centered2 / Smagorinsky: LWP 35.70 g/m²; cloud 99.59%; peak w^2 0.427 m²/s²; Δq_t 0.250 g/kg; E 6.57 mm/s.

Fine / WENO9 / no SGS: LWP 55.48 g/m²; cloud 99.95%; peak w^2 0.472 m²/s²; Δq_t 0.031 g/kg; E 5.34 mm/s.

Fine / WENO5 / no SGS: LWP 57.54 g/m²; cloud 99.90%; peak w^2 0.483 m²/s²; Δq_t 0.007 g/kg; E 5.24 mm/s.

Fine / WENO9 / Smagorinsky: LWP 53.23 g/m²; cloud 99.89%; peak w^2 0.445 m²/s²; Δq_t 0.036 g/kg; E 5.09 mm/s.

Fine / WENO5 / Smagorinsky: LWP 51.71 g/m²; cloud 99.77%; peak w^2 0.455 m²/s²; Δq_t 0.034 g/kg; E 5.48 mm/s.

Fine / Centered2 / Smagorinsky: LWP 44.17 g/m²; cloud 99.98%; peak w^2 0.373 m²/s²; Δq_t 0.154 g/kg; E 6.44 mm/s.

Fourth-hour metrics

Peak w^2 is the maximum of the fourth-hour mean native-face profile; entrainment includes subsidence. These are single-seed results.

Fine-grid averaging-window comparison

Compare the same statistic across averaging windows. Dimensional third moment and skewness differ in normalization as well as the windows used in the two papers. All minima below are over 600–760 m and may occur at different heights. Skewness is the ratio of time-averaged moments.

fine_none_weno9, 2-4 h: minimum w'^3 -0.0497 m³/s³; minimum skewness -0.1503.

fine_none_weno9, 3-4 h: minimum w'^3 -0.0622 m³/s³; minimum skewness -0.1928.

fine_none_weno5, 2-4 h: minimum w'^3 -0.0433 m³/s³; minimum skewness -0.1334.

fine_none_weno5, 3-4 h: minimum w'^3 -0.0372 m³/s³; minimum skewness -0.1109.

fine_smagorinsky_weno9, 2-4 h: minimum w'^3 -0.0419 m³/s³; minimum skewness -0.1394.

fine_smagorinsky_weno9, 3-4 h: minimum w'^3 -0.0304 m³/s³; minimum skewness -0.1026.

fine_smagorinsky_weno5, 2-4 h: minimum w'^3 -0.0163 m³/s³; minimum skewness -0.0538.

fine_smagorinsky_weno5, 3-4 h: minimum w'^3 -0.0160 m³/s³; minimum skewness -0.0521.

fine_smagorinsky_centered2, 2-4 h: minimum w'^3 -0.0236 m³/s³; minimum skewness -0.1206.

fine_smagorinsky_centered2, 3-4 h: minimum w'^3 -0.0272 m³/s³; minimum skewness -0.1309.


How to read these comparisons

All scientific figures are generated in Julia with CairoMakie. Numerical configurations use fixed Okabe-Ito colors, distinct dashes and marker shapes; different resolutions occupy panels with shared scales. The cloud-fraction zoom follows the simulated range, while a separate full-range panel preserves the paper ensemble spread.

Light gray is the paper's model-ensemble minimum–maximum; darker gray is its interquartile range. These are not observational uncertainty bands or validation tolerances. The mean and envelope were digitized from Figure 2 at approximately 0.2–3.9 h. Original pixel coordinates, calibration, extraction settings and an overlay are retained. Subpixel early cloud-fraction bands are left missing. Only rendered cloud-fraction band edges are clipped to the physical 0–1 range; CSV values preserve raw digitization. The zoom panel omits the paper bands to avoid magnifying raster-edge uncertainty.


Paper observations are black markers in profile plots. Measurement error bars were not digitized. Numerical configurations use redundant color, line and symbol cues; resolution panels share scales.

Cloud persistence and turbulence evolution




Gray bands: digitized MODEL spread (Stevens et al. 2005, Fig. 2), not observational uncertainty. Cloud-fraction zoom shows Breeze curves; the full-range row shows the paper envelope. Unresolved early shading is omitted. TKE is resolved only. All plots: Julia/CairoMakie.

Fourth-hour thermodynamic structure




Fourth-hour means combine the two half-hour profile averages ending at 12600 and 14400 s. Black squares: retained paper Fig. 4 observations; measurement error bars were not digitized. Shared axis scales across resolutions.

Fourth-hour vertical-velocity moments




Fourth-hour means on native w faces. Black squares: retained in-situ observations; black open circles: radar (paper Fig. 5). Error bars were not digitized. Colors, line patterns and marker shapes identify numerical configurations; rows identify resolutions.

Fourth-hour scalar transport and buoyancy production




Scalar panels add resolved covariance and explicit SGS flux; buoyancy production is resolved only. Numerical transport is not separately measured, so resolved + SGS is not a complete discrete flux budget.

Cloud boundaries and decoupling



Cloud-base height is conditional on cloudy columns. Inversion height uses the interpolated q_t contour. $\Delta q_t = \text{mean}(100\text{--}200 \text{ m}) - \text{mean}(700\text{--}800 \text{ m})$, using exact geometric layer overlap.

Resolution and numerical choices






Only completed cases appear. LWP bars span two half-hour means, not confidence intervals. Marker RMSE is an unweighted descriptive distance to retained in-situ observations, not an uncertainty-normalized validation score.

Unbounded Centered2 + Smagorinsky: moisture extrema

Coarse: $\Delta x = 80$ m, $\Delta z = 20$ m

Canonical: $\Delta x = 35$ m, $\Delta z = 5$ m







Fine: $\Delta x = 10$ m, $\Delta z = 5$ m



- ··· Minimum q_t (star markers) - - - Maximum q_t ····· Zero total water

Global spatial extrema sampled every 60 s; excursions between samples are not bounded by these curves.
Negative total water is physically inadmissible. Finite output does not establish physical validity; no clipping or WENO fallback was applied.


Canonical grid: Breeze and Pressel et al. (2017)



-- WENO9 / no SGS (99NN) -- WENO9 / Smagorinsky (99MS) -- Centered2 / Smagorinsky (22MS) --- Pressel model
-- WENO5 / no SGS (55NN) -- WENO5 / Smagorinsky (55MS) — Breeze


35 m horizontal / 5 m vertical; same domain. Profiles average 2–4 h. Cloud-fraction axis zoomed. Only completed Breeze cases shown.
 Scheme-family comparison: Cs, Pr, CFL, thermodynamics, surface treatment and moisture handling differ; this is not a replication.

Breeze and Pressel Figure 7: skewness, hours 2–4



Canonical grid: $\Delta x = 35$ m, $\Delta z = 5$ m. Breeze uses the ratio of time-averaged central moments. Pressel normalization/averaging order is unverified; these are model curves, not observations. Top: 0–1000 m. Bottom: magnified 500–800 m. Low-variance levels masked. Model settings differ.

Breeze skewness across resolutions: hours 2–4



Ratio of 2–4 h averaged central moments; not the time average of instantaneous skewness.
Only completed exports shown. Coarse horizontal domain is larger. Variance $\leq 10^{-6} \text{ m}^2/\text{s}^2$ is masked; full-height values are in CSV.

Methods and limitations

Experimental design

Three grids: coarse $96 \times 96 \times 75$ at $80 \times 80 \times 20$ m in a $7.68 \times 7.68 \times 1.5$ km domain; canonical $96 \times 96 \times 300$ at $35 \times 35 \times 5$ m in $3.36 \times 3.36 \times 1.5$ km; fine $336 \times 336 \times 300$ at $10 \times 10 \times 5$ m in the same canonical domain. Fine $dz=5$ m is the stated implementation assumption. Each grid crosses WENO9/WENO5 with no closure/Smagorinsky-Lilly ($Cs=0.16$, $Pr=1$). WENO moisture uses bounds=(0,1) at both orders; order changes apply to all advected variables. Three additional runs apply Centered(order=2) to momentum and all scalars, including moisture, with the same Smagorinsky-Lilly settings: 15 cases total. Centered moisture is unbounded, so comparisons with WENO change scheme family and bounding as well as order. Finite diagnostics and moisture extrema must be checked before interpreting these cases.

Physics and initial conditions

Breeze uses anelastic dynamics, warm-phase saturation adjustment, fixed surface sensible/latent heat fluxes $15/115$ W/m², drag coefficient 0.0011, divergence $D=3.75 \times 10^{-6}$ s⁻¹, geostrophic winds (7,-5.5) m/s and the RF01 idealized longwave radiation. Initial inversion is 840 m, $\theta_l=289$ K and $q_t=9$ g/kg below, $\theta_l=297.5+(z-840)^{(1/3)}$ K and $q_t=1.5$ g/kg above. No drizzle or droplet sedimentation. All runs use Float32, seed 123 and four simulated hours. The fixed-flux moisture conversion uses $L_v=2.47 \times 10^6$ J/kg; Breeze retains its default thermodynamic latent-heat law with dry-air $c_p=1015$ J/(kg K), which differs from a strictly constant-mixture- c_p implementation.

Averaging and grid locations

Profiles are averages of instantaneous horizontally reduced statistics. The fourth hour combines the equal-duration bins ending at 12,600 and 14,400 s; it excludes the bin ending at 10,800 s. Scalar means are on centers; w variance and third central moment remain on native vertical faces. The third moment has units m³/s³, not normalized skewness. Series means use trapezoidal integration of 60 s samples between 10,800 and 14,400 s. Two half-hour LWP means give a descriptive temporal spread, not a confidence interval.

Methods and limitations

Cloud and entrainment definitions

Column cloud fraction uses $q_l > 10^{-6}$ kg/kg in any cell. Cloud-base means exclude clear columns. Inversion height uses the first downward $q_t = 8$ g/kg crossing with linear interpolation; its valid fraction is checked. The $\theta_{\text{ref}} = 295$ K contour and the nearest-cell height used internally by radiation are retained separately. Fourth-hour entrainment is $[z_i(4h) - z_i(3h)]/3600 + D \cdot \text{mean}(z_i)$, including subsidence. Decoupling is mean q_t over 100–200 m minus mean q_t over 700–800 m, with exact layer-overlap weights.

Reference data and comparison limits

Stevens et al. (2005), DOI 10.1175/MWR2930.1, supplies the RF01 comparison. Figure 2 is a model ensemble mean, not observations or an acceptance band; the paper has no observed LWP time series. Observed cloud cover exceeds 99%. Figure 4 and 5 markers are partial digitizations, with source pixels retained and error bars not extracted. Marker RMSE is an unweighted descriptive distance to retained in-situ points; it is not a likelihood, uncertainty-normalized score, or formal validation criterion. Radar markers are plotted separately. Resolved TKE is plotted against the historical ensemble for context; a prognostic SGS-TKE contribution is unavailable for these closures.

What this design can establish

The coarse domain has 5.22 times the canonical horizontal area, so coarse-versus-canonical differences combine changes in horizontal spacing, vertical spacing and domain size. Fine-versus-canonical holds domain and vertical spacing fixed. Closure/order comparisons share a seed within a grid, but one seed gives no ensemble uncertainty and different grids do not share identical perturbation fields. UCLA-0 in the paper disables scalar SGS mixing while retaining momentum treatment; it is not identical to Breeze closure=nothing. Neither matching LWP nor agreement with one fine simulation establishes convergence.

Methods and limitations

Numerical and output audit

The original WENO runs use an immutable source snapshot with per-file hashes and a pinned Julia 1.12.6 workspace. The Centered2 extension uses a separate versioned runner and provenance while preserving the original source. GPU smoke runs checked all grids and the diagnostic output path before production. Exports require completed jobs, end-of-run sentinels, finite numeric records, exactly nine profile times and 241 series times, 46 profile variables and 32 series variables. Native face/center coordinates and fixed fluxes are checked again locally. Replicated budget arrays, where present, are horizontally averaged by the exporter and their original shape/spread recorded. The TKE residual is a budget-closure residual, not a measured numerical dissipation.

Reproducibility and current run state

All active processing and figures use Julia. CairoMakie version 0.15.14; Julia 1.12.7. The companion Markdown/HTML report links to all CSVs, source hashes, digitization provenance and exact rebuild commands. Original restart checkpoints and JLD2 files remain on pcluster.

coarse_none_weno9: COMPLETED

coarse_none_weno5: COMPLETED

coarse_smagorinsky_weno9: COMPLETED

coarse_smagorinsky_weno5: COMPLETED

canonical_none_weno9: COMPLETED

canonical_none_weno5: COMPLETED

canonical_smagorinsky_weno9: COMPLETED

canonical_smagorinsky_weno5: COMPLETED

fine_none_weno9: COMPLETED

fine_none_weno5: COMPLETED

fine_smagorinsky_weno9: COMPLETED

fine_smagorinsky_weno5: COMPLETED

coarse_smagorinsky_centered2: COMPLETED

canonical_smagorinsky_centered2: COMPLETED

fine_smagorinsky_centered2: COMPLETED

Reference: Stevens et al. (2005), Monthly Weather Review 133, 1443–1462. DOI 10.1175/MWR2930.1.

GABLS1: setup and results

Part II: GABLS1

12/15 completed exports admitted to the scientific comparison. Checked 2026-09-20T01:47:54.123Z.

Grids: 12.5, 6.25, 3.125, 2, and 1 m isotropic; 32^3 , 64^3 , 128^3 , 200^3 , and 400^3 cells, all in the same 400 m cube. Small grids run first.

Each grid has exactly three configurations: WENO9 with no interior closure, WENO5 with no interior closure, and WENO9 with Smagorinsky. Smagorinsky $C_s=0.16$, $Pr=1$; paired seed 123 within each grid. These are controlled Breeze settings, not a reproduction of one named historical LES member.

Forcing: $U_g = 8$ m/s, $V_g = 0$, $f = 1.39e-4/s$. Initial potential temperature 265 K below 100 m and $265+0.01(z-100)$ K above. Initial temperature noise is uniform with 0.1 K peak-to-peak range below 50 m. Surface $\theta_{s,ref}=265-0.25t_{hours}$ K; $\theta_{ref}=263.5$ K is a separate buoyancy reference.

Wall: MOST exchange with $\kappa = 0.4$, $\beta_m = 4.8$, $\beta_h = 7.8$, roughness = 0.1 m for heat and momentum. Impermeable boundaries, free-slip upper lid. Vertical-velocity damping is zero below 300 m and ramps linearly to a 60 s relaxation time at 400 m. MOST uses a minimum wind of 0.01 m/s and caps z/L at 10, with a neutral fallback for nonpositive bulk Richardson number; their activity is diagnosed. Exact implementation, damping and timestep choices are recorded in production provenance.

Diagnostics: 541 instantaneous series records at 60 s intervals; 19 profile times with preceding 30 min means after instantaneous t_0 . Compare 7-8 h and 8-9 h. Final-hour profiles combine bins ending at 30600 and 32400 s; penultimate-hour bins end at 27000 and 28800 s.

Compare mean $u/v/\theta$, jet speed/height/turning, total-stress boundary-layer height, surface heat flux/ u_{star}/L , native-face $w_2/w_3/skewness$, resolved/SGS flux partitions and available TKE-budget terms. Missing SGS energy or budget terms remain unavailable; a residual is not numerical dissipation.

Beare et al. (2006), Boundary-Layer Meteorology 118, 247-272, doi:10.1007/s10546-004-2820-6. Official reduced data archive contains 423 files. Native coordinates and missing values are retained; some NERSC headers give 474-534 min for nominal 8-9 h, so exact window matching is not universal.

GABLS1: setup and results

Cluster agents own implementation, scheduling and audited exports. Desktop monitoring is scheduled every 10 min to collect newly completed cases and rebuild this document. Actual submitted jobs and completed-export evidence are recorded in the status files, separate from the authorized matrix.

Production submission records: 15/15 cases have assigned Slurm jobs. Submitted is distinct from running or scientifically complete. The live queue snapshot below also includes any smoke tests.

Implementation choices: Breeze uses anelastic dynamics and Float32 arithmetic; the reference archive is predominantly Boussinesq. Base pressure is 100000 Pa so the prescribed surface temperature equals surface potential temperature; the corresponding reference density differs from the specification's rounded 1.3223 kg/m³ by about 8 parts per million. Smagorinsky-Lilly uses $C_s=0.16$, $C_b=1$, $Pr=1$.

At $\Delta=12.5$ m, WENO9/Smagorinsky has strongly suppressed resolved turbulence: peak final-hour w_2 is $1.12e-06$ m²/s², $1.21e-05$ times the no-closure value. Maximum magnitudes of the final-hour mean resolved and SGS momentum-flux vectors are $1.34e-07$ and 0.0415 m²/s². Its h is 140.6 m with only 0.242 m standard deviation across final-hour 60 s samples. Similar h or mean wind can therefore coexist with very different resolved turbulence and transport partition; h alone cannot establish fidelity. Skewness in this weak-variance regime must not be interpreted as evidence of vigorous turbulent motions.

At $\Delta=12.5$ m without interior closure, WENO9 versus WENO5 gives final-hour h 195.6 versus 136.1 m and peak resolved w_2 0.0927 versus 0.0574 m²/s². This is a scheme sensitivity at fixed grid and seed, not evidence of resolution convergence.

At $\Delta=6.25$ m without interior closure, WENO9 versus WENO5 gives final-hour h 187.8 versus 150.9 m and peak resolved w_2 0.0976 versus 0.0656 m²/s². This is a scheme sensitivity at fixed grid and seed, not evidence of resolution convergence.

At $\Delta=3.12$ m without interior closure, WENO9 versus WENO5 gives final-hour h 184.4 versus 162.9 m and peak resolved w_2 0.1008 versus 0.0696 m²/s². This is a scheme sensitivity at fixed grid and seed, not evidence of resolution convergence.

At $\Delta=2$ m without interior closure, WENO9 versus WENO5 gives final-hour h 178.5 versus 159.2 m and peak resolved w_2 0.0993 versus 0.0690 m²/s². This is a scheme sensitivity at fixed grid and seed, not evidence of resolution convergence.

GABLS1: setup and results

n032_weno9_none: peak final-hour resolved w_2 ($0.0927 \text{ m}^2/\text{s}^2$) exceeds the largest archived model peak at the same spacing ($0.0602 \text{ m}^2/\text{s}^2$). Model spread is a comparison baseline, not observational truth; stronger resolved turbulence alone does not establish better fidelity.

n064_weno9_none: peak final-hour resolved w_2 ($0.0976 \text{ m}^2/\text{s}^2$) exceeds the largest archived model peak at the same spacing ($0.0779 \text{ m}^2/\text{s}^2$). Model spread is a comparison baseline, not observational truth; stronger resolved turbulence alone does not establish better fidelity.

n128_weno9_none: peak final-hour resolved w_2 ($0.1008 \text{ m}^2/\text{s}^2$) exceeds the largest archived model peak at the same spacing ($0.0904 \text{ m}^2/\text{s}^2$). Model spread is a comparison baseline, not observational truth; stronger resolved turbulence alone does not establish better fidelity.

Resolution comparison for WENO9/none: reducing spacing from 12.5 to 2 m changes final-hour h from 195.6 to 178.5 m and peak resolved w_2 from 0.0927 to $0.0993 \text{ m}^2/\text{s}^2$. Surface heat flux changes from -15.52 to -11.87 W/m^2 . These compare the coarsest and finest completed grids for this configuration; similarity of selected quantities alone is not a convergence or fidelity test.

Resolution comparison for WENO5/none: reducing spacing from 12.5 to 2 m changes final-hour h from 136.1 to 159.2 m and peak resolved w_2 from 0.05744 to $0.06902 \text{ m}^2/\text{s}^2$. Surface heat flux changes from -11.79 to -10.04 W/m^2 . These compare the coarsest and finest completed grids for this configuration; similarity of selected quantities alone is not a convergence or fidelity test.

Resolution comparison for WENO9/smagorinsky: reducing spacing from 12.5 to 2 m changes final-hour h from 140.6 to 188.9 m and peak resolved w_2 from $1.12\text{e-}06$ to $0.09298 \text{ m}^2/\text{s}^2$. Surface heat flux changes from -9.17 to -12.22 W/m^2 . These compare the coarsest and finest completed grids for this configuration; similarity of selected quantities alone is not a convergence or fidelity test.

For WENO9/Smagorinsky, the finest completed grid has $8.3\text{e}+04$ times the coarsest-grid peak final-hour resolved w_2 . At $\Delta=2$ m, maximum magnitudes of final-hour mean resolved and SGS momentum-flux vectors are 0.06499 and $0.04146 \text{ m}^2/\text{s}^2$; these maxima can occur at different heights and should not be treated as fractions of a common total. Final-hour h sample standard deviation is 4.63 m, and mean h changed by 3.31 m between 7–8 and 8–9 h. Read the variance evolution and flux profiles together: this resolution sensitivity does not establish statistical stationarity or identify a unique cause.

GABLS1: setup and results

Final-hour values are descriptive results, not a fidelity ranking. Resolution comparisons must keep the scheme and closure fixed; archive ensemble membership changes with resolution.

The archive reports skewness directly; its averaging convention is not assumed identical to our ratio of time-averaged central moments. No archive third moment is reconstructed from independently averaged variance and skewness. Both Breeze skewness plots mask levels with final-hour variance at or below $1\text{e-}8 \text{ m}^2/\text{s}^2$; the raw diagnostics remain available.

A fixed 1.0 m archive ensemble median (black dotted line) is repeated across resolution panels wherever a matching quantity is available, alongside same-grid archive spread. The 1.0 m archive has two models; pointwise member counts and source files are saved in `reference_data/fixed_1m_medians.json`. Linear totals and resolved TKE are computed per model before taking medians. No matching reference is fabricated for `w3`, time-mean stress magnitude, unmatched budget terms, wall-regime fractions, peak-variance time series or jet time series. The fixed reference aids resolution comparisons but is not an exact solution.

The 1.0 m IMUK archive contains conspicuously large upper-level SGS values, confirmed in the original files: SGS `uw` spans -0.872 to $0.869 \text{ m}^2/\text{s}^2$ and SGS shear production spans -0.283 to $0.793 \text{ m}^2/\text{s}^3$. Tiny values approaching $1\text{e-}100$ abruptly become order-one values above roughly 330 m. This pattern suggests an archive exponent-format problem, but the intended values have not been established. Raw values and resulting median curves are retained, with affected flux/budget panels flagged. Do not interpret those upper-level excursions as a convergence target.

Corrected production (v2) replaces the original campaign after a diagnostic index restriction truncated total-flux/stress profiles and exposed a boundary-layer-height contour-index error. The corrected source preserves the original physics and all three numerical configurations. Expanded CPU checks passed 278 kernel and 426 writer checks; GPU checks passed 268 kernel and 422 writer checks, including every raw profile extent. All 15 cases were resubmitted as arrays 7099-7103, beginning with 32³. Original outputs remain excluded from scientific comparisons and preserved as failure evidence. The collector and single cluster exporter now target v2.

GABLS1: setup and results

The original 400³ WENO9/no-closure and WENO9/Smagorinsky attempts failed during startup and remain excluded. Their fixed initial timestep of 0.5 s implied an advective CFL of 4 for 8 m/s wind and 1 m cells. The separate startup_dt_v1 snapshot changes only initial dt to $\min(0.5, 0.5 \cdot dx/8)$, or 0.0625 s at 1 m. Both variants passed 72 s GPU tests with finite saved bounds: 852/832 iterations and final actual CFL 0.700088/0.700007. The unchanged wizard gave an effective first dt of 0.06875 s. Replacement jobs 7120_1 (WENO9/no closure) and 7120_2 (WENO9/Smagorinsky) were submitted with a one-job array limit; 7120_1 is running and 7120_2 queued at this snapshot. Original WENO5 job 7103_2 continues without duplication. Original failed source/logs/outputs are preserved in attempt history, and replacements require their own job identity, revision and source-hash audits. Passing the short startup tests supports the repair but does not prove full-run stability, uniquely establish the failure cause, or establish that coarser results are independent of startup settings.

n032_weno9_none: final-hour $h = 195.6$ m; $u_{star} = 0.281$ m/s; surface heat flux = -15.52 W/m²; peak mean $w_2 = 0.0927$ m²/s². Between 7–8 and 8–9 h, h changed by -3.7 m and the w_2 profile changed by 0.00438 m²/s² RMS. Simulation wall time 529.7 s (40114 iterations), excluding preceding compilation/startup.

n032_weno5_none: final-hour $h = 136.1$ m; $u_{star} = 0.242$ m/s; surface heat flux = -11.79 W/m²; peak mean $w_2 = 0.05744$ m²/s². Between 7–8 and 8–9 h, h changed by -2.3 m and the w_2 profile changed by 0.00379 m²/s² RMS. Simulation wall time 505.5 s (39025 iterations), excluding preceding compilation/startup.

n032_weno9_smagorinsky: final-hour $h = 140.6$ m; $u_{star} = 0.223$ m/s; surface heat flux = -9.17 W/m²; peak mean $w_2 = 1.12e-06$ m²/s². Between 7–8 and 8–9 h, h changed by 0.1 m and the w_2 profile changed by $4.74e-08$ m²/s² RMS. Simulation wall time 473.4 s (35312 iterations), excluding preceding compilation/startup.

n064_weno9_none: final-hour $h = 187.8$ m; $u_{star} = 0.279$ m/s; surface heat flux = -15.24 W/m²; peak mean $w_2 = 0.09765$ m²/s². Between 7–8 and 8–9 h, h changed by -3.3 m and the w_2 profile changed by 0.00154 m²/s² RMS. Simulation wall time 1050.0 s (80959 iterations), excluding preceding compilation/startup.

n064_weno5_none: final-hour $h = 150.9$ m; $u_{star} = 0.248$ m/s; surface heat flux = -12.01 W/m²; peak mean $w_2 = 0.06558$ m²/s². Between 7–8 and 8–9 h, h changed by -10.9 m and the w_2 profile changed by 0.00248 m²/s² RMS. Simulation wall time 1029.6 s (79613 iterations), excluding preceding compilation/startup.

GABLS1: setup and results

n064_weno9_smagorinsky: final-hour $h = 149.7$ m; $u_{star} = 0.234$ m/s; surface heat flux = -10.16 W/m²; peak mean $w2 = 0.0315$ m²/s². Between 7–8 and 8–9 h, h changed by -21.0 m and the $w2$ profile changed by 0.00986 m²/s² RMS. Simulation wall time 1020.2 s (79350 iterations), excluding preceding compilation/startup.

n128_weno9_none: final-hour $h = 184.4$ m; $u_{star} = 0.270$ m/s; surface heat flux = -13.56 W/m²; peak mean $w2 = 0.1008$ m²/s². Between 7–8 and 8–9 h, h changed by -1.3 m and the $w2$ profile changed by 0.000826 m²/s² RMS. Simulation wall time 4156.8 s (162036 iterations), excluding preceding compilation/startup.

n128_weno5_none: final-hour $h = 162.9$ m; $u_{star} = 0.245$ m/s; surface heat flux = -11.12 W/m²; peak mean $w2 = 0.0696$ m²/s². Between 7–8 and 8–9 h, h changed by 7.9 m and the $w2$ profile changed by 0.00189 m²/s² RMS. Simulation wall time 4058.5 s (160877 iterations), excluding preceding compilation/startup.

n128_weno9_smagorinsky: final-hour $h = 188.1$ m; $u_{star} = 0.277$ m/s; surface heat flux = -13.37 W/m²; peak mean $w2 = 0.08562$ m²/s². Between 7–8 and 8–9 h, h changed by 3.9 m and the $w2$ profile changed by 0.00112 m²/s² RMS. Simulation wall time 4172.6 s (161379 iterations), excluding preceding compilation/startup.

n200_weno9_none: final-hour $h = 178.5$ m; $u_{star} = 0.261$ m/s; surface heat flux = -11.87 W/m²; peak mean $w2 = 0.0993$ m²/s². Between 7–8 and 8–9 h, h changed by 3.3 m and the $w2$ profile changed by 0.000919 m²/s² RMS. Simulation wall time 9396.4 s (253504 iterations), excluding preceding compilation/startup.

n200_weno5_none: final-hour $h = 159.2$ m; $u_{star} = 0.240$ m/s; surface heat flux = -10.04 W/m²; peak mean $w2 = 0.06902$ m²/s². Between 7–8 and 8–9 h, h changed by 1.8 m and the $w2$ profile changed by 0.000677 m²/s² RMS. Simulation wall time 8387.3 s (252093 iterations), excluding preceding compilation/startup.

n200_weno9_smagorinsky: final-hour $h = 188.9$ m; $u_{star} = 0.274$ m/s; surface heat flux = -12.22 W/m²; peak mean $w2 = 0.09298$ m²/s². Between 7–8 and 8–9 h, h changed by 3.3 m and the $w2$ profile changed by 0.00138 m²/s² RMS. Simulation wall time 10660.3 s (252948 iterations), excluding preceding compilation/startup.


Current queue snapshot (job|state|elapsed|start|node or pending reason):

7120_2|PENDING|0:00|N/A|(JobArrayTaskLimit)

7120_1|RUNNING|3:13|2026-09-20T01:44:41|gpu-prod-st-gpu-prod-1


7103_2|RUNNING|2:16:48|2026-09-19T23:31:06|gpu-prod-st-gpu-prod-2

GABLS1 reference archive: final-hour model spread



Archive: nominal final-hour profiles, including some NERSC windows of 474–534 min. Breeze curves, when present: 8–9 h means.
 Native heights interpolated without extrapolation; model membership varies with height and grid. The spread is not observational uncertainty.
 Black dots: fixed 1.0 m ensemble median; interpolation precedes the median. Unavailable references are labeled.


GABLS1 reference archive: evolution



Same-grid archive min-max (≥ 2 members)
 Same-grid archive median
 1.0 m archive median (fixed reference)


Archive series interpolated to common times without extrapolation; missing values omitted.
 Membership varies with grid and time. Black dots: the same 1.0 m ensemble median in every row. Heat flux is upward positive.

GABLS1: mean wind and temperature




Archive: nominal final-hour profiles, including some NERSC windows of 474–534 min. Breeze curves, when present: 8–9 h means.
Native heights interpolated without extrapolation; model membership varies with height and grid. The spread is not observational uncertainty.
Black dots: fixed 1.0 m ensemble median; interpolation precedes the median. Unavailable references are labeled.

GABLS1: vertical velocity moments




Archive: nominal final-hour profiles, including some NERSC windows of 474–534 min. Breeze curves, when present: 8–9 h means.
Native heights interpolated without extrapolation; model membership varies with height and grid. The spread is not observational uncertainty.
Black dots: fixed 1.0 m ensemble median; interpolation precedes the median. Unavailable references are labeled.

GABLS1: streamwise momentum transport




Archive: nominal final-hour profiles, including some NERSC windows of 474–534 min. Breeze curves, when present: 8–9 h means.
Native heights interpolated without extrapolation; model membership varies with height and grid. The spread is not observational uncertainty.
Black dots: fixed 1.0 m ensemble median; interpolation precedes the median. Unavailable references are labeled.
Archive caution: 1.0 m IMUK SGS data contain upper-level anomalies; retained as supplied and flagged in the report.

GABLS1: vertical heat transport




Archive: nominal final-hour profiles, including some NERSC windows of 474–534 min. Breeze curves, when present: 8–9 h means.
Native heights interpolated without extrapolation; model membership varies with height and grid. The spread is not observational uncertainty.
Black dots: fixed 1.0 m ensemble median; interpolation precedes the median. Unavailable references are labeled.
Archive caution: 1.0 m IMUK SGS data contain upper-level anomalies; retained as supplied and flagged in the report.

GABLS1: turbulence intensity, asymmetry and stress




Archive: nominal final-hour profiles, including some NERSC windows of 474–534 min. Breeze curves, when present: 8–9 h means.
Native heights interpolated without extrapolation; model membership varies with height and grid. The spread is not observational uncertainty.
Black dots: fixed 1.0 m ensemble median; interpolation precedes the median. Unavailable references are labeled.

GABLS1: selected TKE budget terms




Breeze: 8–9 h means; only fully completed and audited exports. Shown budget terms do not form a closed budget.
A budget residual must not be interpreted as numerical dissipation; unavailable terms are not replaced by zero.
Black dots: fixed 1.0 m ensemble median; interpolation precedes the median. Unavailable references are labeled.
Archive caution: 1.0 m IMUK SGS data contain upper-level anomalies; retained as supplied and flagged in the report.

GABLS1: crosswind momentum transport




Archive: nominal final-hour profiles, including some NERSC windows of 474–534 min. Breeze curves, when present: 8–9 h means.
Native heights interpolated without extrapolation; model membership varies with height and grid. The spread is not observational uncertainty.
Black dots: fixed 1.0 m ensemble median; interpolation precedes the median. Unavailable references are labeled.
Archive caution: 1.0 m IMUK SGS data contain upper-level anomalies; retained as supplied and flagged in the report.

GABLS1: wall-model regimes and resolved turbulence




Instantaneous 60 s samples. The wall law caps ζ at 10 and uses neutral exchange for nonpositive bulk Richardson number. These fractions identify when wall-model limits participate in the result. No matching archive regime or peak-variance series is available.

GABLS1: resolved vertical-velocity variance



Logarithmic variance axes distinguish weak resolved motions from sustained turbulence. Left: instantaneous profile maxima; right: 8–9 h means. Black dots: fixed 1.0 m median (right only). No archived peak-variance time series. Nonpositive values omitted without a substituted floor.


GABLS1: evolution



— WENO9 / no SGS - - - WENO9 / Smagorinsky - - - Same-grid archive median
 - - - WENO5 / no SGS Same-grid archive min-max (≥ 2 members) ···· 1.0 m archive median (fixed reference)

Breeze: instantaneous 60 s samples. Archive series interpolated to common times without extrapolation; model membership varies.
 h = height of 5% total surface stress / 0.95; invalid crossings omitted. Black dots: fixed 1.0 m median. Heat flux is upward positive.

GABLS1: low-level jet evolution



Instantaneous 60 s samples of the maximum horizontal-mean wind speed. Direction is atan(v, u) relative to eastward geostrophic wind. Invalid jet diagnostics are omitted; a changing peak height can jump between adjacent grid levels. No archived jet time series is available.

Part III: GABLS3 preparation

GABLS3: observational evaluation and case preparation

Status: preparation, with no GABLS3 production results yet. All DYCOMS and GABLS1 material remains in this master report. GABLS1's fixed 1 m ensemble reference is retained in its own comparison.

Published LES case

The original GABLS3 LES experiment covers 00-09 UTC on 2 July 2006 at Cabauw. It uses an 800 m cube with a common 6.25 m grid. Initial profiles combine tower, wind-profiler and sounding measurements. Geostrophic wind and advective tendencies vary in time and height. Prescribed potential temperature and humidity at 0.25 m replace interactive land-surface and radiation schemes; aerodynamic roughness is 0.15 m. Basu, Holtslag and Bosveld (2012), sections 2-3.

For 03-04 UTC, the ensemble reproduced jet height, strength, directional shear and fluxes reasonably well, but near-surface shear was deficient. A 1 m simulation mixed less and produced a stronger jet. Observed submesoscale variability and a turbulence burst were missed. Published LES variances are resolved, whereas observations include total variance. The ensemble percentiles pool models and five-minute samples; tower whiskers show hourly minima/maxima, not confidence intervals. Basu et al. (2012), sections 4-5 and figures 1-4.

The related 24-hour SCM experiment begins at noon on 1 July and couples the atmosphere to land-surface and radiation schemes. Its input tables must not silently replace the LES specification. Bosveld et al. (2014), case selection and setup.

Breeze experiment under preparation

Our proposed resolution comparison keeps the 800 m cube fixed: 12.5 m (64 cubed), 6.25 m (128 cubed), and 3.125 m (256 cubed). These nine proposed cases use exactly the user's three configurations: WENO9 without interior closure, WENO5 without interior closure, and WENO9 with Smagorinsky. The added coarse and fine grids are our experiment design, not a claim about the historical intercomparison matrix. Case preparation and validation precede production scheduling; no GABLS3 production run is claimed here.

We will verify the initial profiles, forcing tables, UTC origin, units, moisture convention, roughness lengths, near-surface reference height, wall stability functions through sunrise, damping and perturbations before freezing the setup. Surface exchange remains active in runs without an interior closure. Unresolved parameters stay explicit; GABLS1 defaults cannot substitute for missing GABLS3 inputs.

Part III: GABLS3 preparation

The comparison will retain observed points and their sampling ranges, published ensemble ranges, and any individual high-resolution LES as distinct references. A single 1 m GABLS3 run is not the two-model GABLS1 median and is not observational truth. Resolution agreement and agreement with observations will be assessed separately.

The author's revised instructions specify instantaneous plane-averaged profiles every five minutes, and surface/point time series every ten seconds. We will preserve that sampling, rather than applying GABLS1 half-hour output averaging. The 0.25 m scalar reference height is distinct from momentum roughness; the wall treatment must handle both stable nighttime and unstable morning conditions. Revised participant specification, archived in the author's repository.

Planned output includes mean u/v /potential temperature/humidity, velocity variances, vertical-velocity third moment and skewness, resolved/SGS/total momentum, heat and moisture fluxes, available TKE terms, jet height and speed, and surface fluxes, friction velocity and Obukhov length. We will compare the paper's 03-04 UTC interval and the complete night-to-morning evolution. Missing observed third moments or SGS quantities will remain unavailable. Numerical budget residuals will not be labeled dissipation without a closed budget.

Preparation checks cover forcing interpolation and jumps, units, profile shapes, surface fluxes and time stepping. GPU smoke tests retain the two-job limit and GABLS1 priority. Scientific admission requires source/input hashes, completion evidence and audited exports.

Persistent evaluation repository

Created as a private repository: `glwagner/BreezeEvaluation.jl`. Migration, GABLS3 preparation and scientific completion are tracked separately. The combined report and all plots use Julia.

The repository will retain DYCOMS and GABLS1 case definitions, pinned dependencies, provenance, audited results, reference data, plotting/report code, figures and this report. Frozen production, failed attempts and raw fields remain on `pcluster` with checksummed manifests. Original simulation identities and hashes are preserved. Third-party papers are linked.

All analysis and plots use Julia. Existing completed experiments and running GABLS1 jobs continue unchanged during migration. Future evaluations will live in `BreezeEvaluation.jl`, with their Breeze dependency revision pinned explicitly.